AI-Powered Incident Triage Copilot

Enhancing BFSI System Reliability with GenAI

Hackathon 2025

Problem Statement

Our TCS BFSI Client handles large volume of Production incidents/Tickets

- Engineers and Site Reliability Teams spend significant time analyzing logs, identifying root causes, and coordinating fixes.
- High MTTA/MTTR impacts business KPIs and customer experience.
- Lack of a centralized AI-driven triaging system delays incident resolution.

Current Incident Triaging Procedure

SRE/SRC teams manually analyze alerts, logs and ETL Job failures from monitoring tools like Dynatrace and Logscale.

Frequent switching between BigPanda, Splunk, CA7 and confluence runbooks to investigate issues.

Root cause Analysis takes time due to scattered knowledge and repetitive tasks.

Collaboration between teams is slow, mostly through emails or manual updates.

Results in high MTTA and High MTTR.

Proposed Solution

When BigPanda creates a new incident, it pushes payloads to the LLM-powered FastAPI Service.

Use LLM+RAG Powered Analysis to monitor Dynatrace metrics.

Detect failed jobs and provide dependency chain for the job and it’s downstream batch.

LLM can also be trained on Confluence Data to retrieve context specific resolution via RAG.

Integration with teams/Slack which provides adaptive cards in-thread which would provide RCA, step by step fixes, validation checks and run-book links.

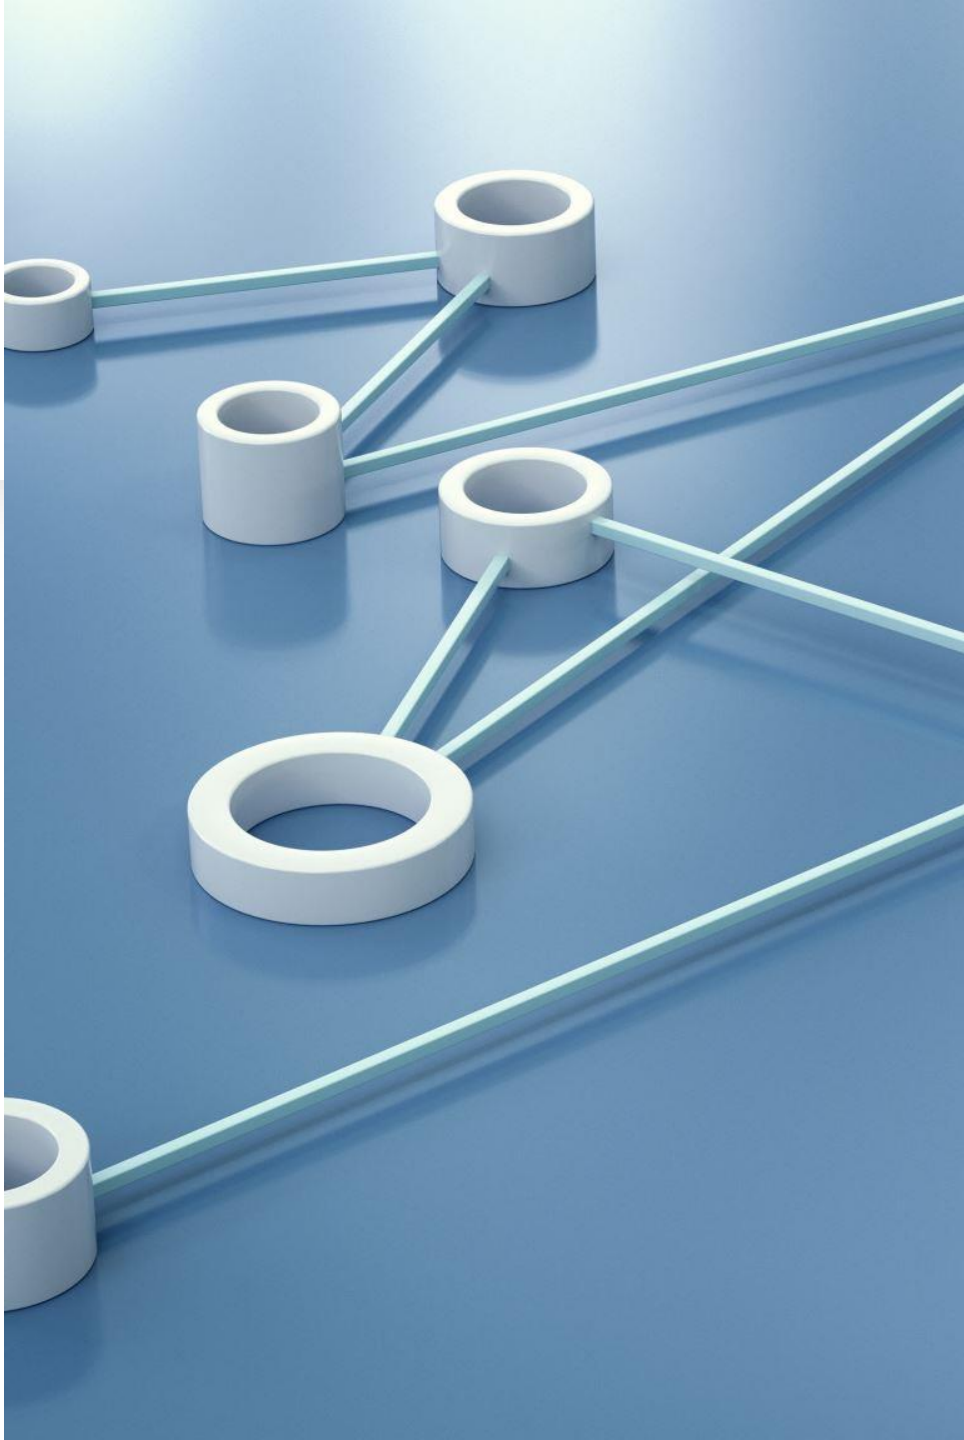

Solution Architecture

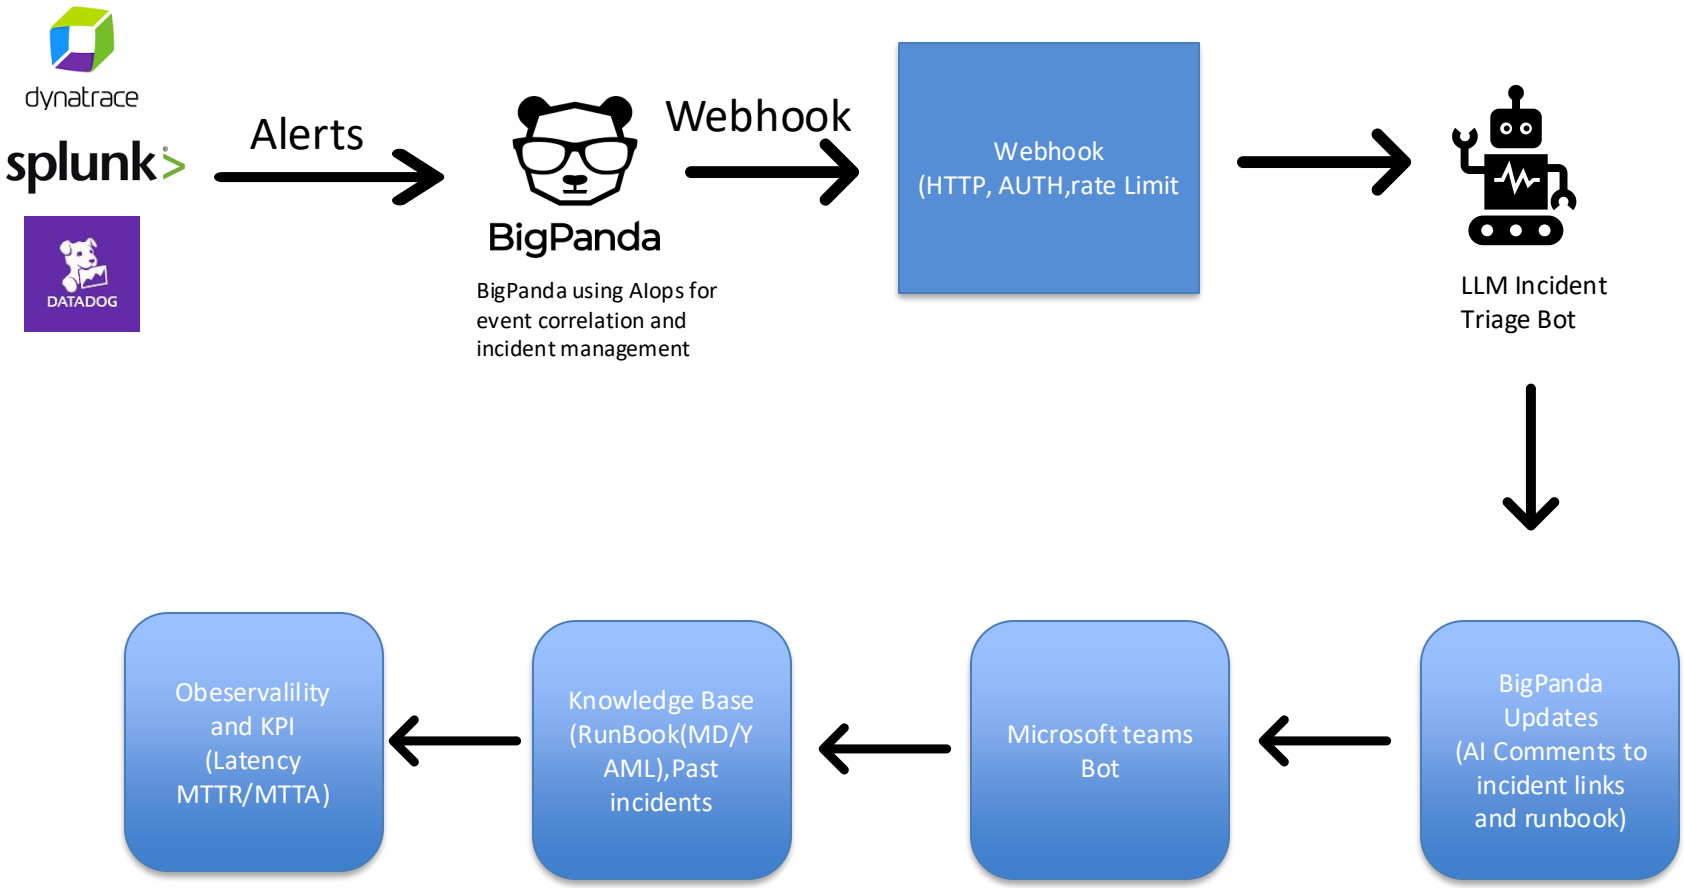

Expected Impact

Reduce	Reduce MTTA from ~15 mins → ~5 mins.
Reduce	Reduce MTTR from ~45 mins → ~20 mins.
Decrease	Decrease manual analysis of repetitive alerts by ~60%.
Improve	Improve engineer productivity and reduce on-call fatigue.
Enhance	Enhance system reliability and customer experience in BFSI.